

Name: Bhavik Pahuja

UID: 24BAI70522

Course: Bachelor of CSE (AI) – 2nd Year

Subject: Database Management Systems

Experiment 2: Advanced Data Aggregation and Filtering

1. Aim of the Session

The aim of this practical is to implement and analyze Group Functions and Conditional Filtering in SQL. The session focuses on using GROUP BY, HAVING, and ORDER BY clauses to extract meaningful insights from an employee dataset.

2. Objective of the Session

By completing this practical, I have achieved the following:

- Developed a schema for employee management using appropriate data types like NUMERIC and DATE.
- Mastered the use of Aggregate Functions (specifically AVG) to perform calculations on data groups.
- Learned to differentiate between the WHERE clause (row-level filtering) and the HAVING clause (group-level filtering).
- Gained proficiency in sorting aggregated results using the ORDER BY clause.

3. Practical / Experiment Steps

The following implementation tasks were completed:

1. Schema Definition: Created the employee table with constraints and precise numeric scaling for salaries.
2. Data Population: Inserted diverse records representing various departments (IT, HR, Sales, Finance) and salary ranges.
3. Basic Aggregation: Calculated the average salary per department using the GROUP BY clause.
4. Advanced Filtering: Applied the HAVING clause to filter out departments where the average salary did not meet a specific threshold.

5. Complex Querying: Combined WHERE, GROUP BY, HAVING, and ORDER BY into a single query to refine results based on individual salaries and group averages simultaneously.

4. Procedure of the Practical

The experiment was conducted following these sequential steps:

1. System Initialization: Logged into the PostgreSQL environment via pgAdmin 4 using localhost as the host server.
2. Table Construction: Executed the CREATE TABLE command to define the structure for the employee dataset.
3. Data Insertion: Ran multiple INSERT statements to populate the table with the provided employee data.
4. Initial Verification: Used SELECT * to confirm that all employee records were correctly stored and formatted.
5. Group Analysis: Executed a GROUP BY query to observe the distribution of average salaries across different departments.
6. Applying Group Filters: Integrated the HAVING clause to restrict the output to high-paying departments (Average > 30,000).
7. Final Refinement: Executed a comprehensive query that filtered individual employees (Salary > 20,000), grouped them by department, and sorted the results in descending order.
8. Output Recording: Captured screenshots of the query results and saved the final SQL script for documentation.

5. I/O Analysis (Input / Output Analysis)

Input Queries

SQL

-- Table creation

```
CREATE TABLE employee(
```

```
    emp_id NUMERIC PRIMARY KEY,
```

```
    emp_name VARCHAR(50),
```

```
    department VARCHAR(50),
```

```
salary NUMERIC(10,2),  
joining_date DATE  
);
```

-- Advanced Aggregation Query

```
SELECT department, AVG(salary) AS avg_salary  
FROM employee  
WHERE salary > 20000  
GROUP BY department  
HAVING AVG(salary) > 30000  
ORDER BY avg_salary DESC;
```

Output Details

- Aggregate Results: The system successfully grouped employees by department.
- Filtering Logic: The WHERE clause correctly excluded employees with salaries under 20,000 (like Sara and Vikram) before calculating averages.
- Group Filtering: The HAVING clause ensured only departments with an average salary exceeding 30,000 were displayed in the final output.
- Sorting: The ORDER BY clause successfully sorted the final results from highest to lowest average salary.

32 ▶ Run | +Tab | JSON | ☰Select | ☱Ask AI

✓ 33 SELECT DEPARTMENT, AVG(SALARY)::NUMERIC(10,2) AS AVG_SAL

34 FROM EMPLOYEE

35 GROUP BY DEPARTMENT

36 ORDER BY AVG(SALARY) DESC; 19ms You, now • Uncommitted changes

EMPLOYEE X

Search Results Export Cost: 26ms Total 3

	department	avg_sal
>	IT	27000.00
>	HR	23000.00
>	Finance	22250.00

27 ▶ Run | +Tab | JSON | ☰Select | ☱Ask AI

✓ 28 SELECT DEPARTMENT, AVG(SALARY)::NUMERIC(10,2) AS AVG_SAL

29 FROM EMPLOYEE

30 GROUP BY DEPARTMENT

31 HAVING AVG(SALARY)>30000; 4ms

32 ▶ Run | +Tab | JSON | ☰Select | ☱Ask AI

33 SELECT DEPARTMENT, AVG(SALARY)::NUMERIC(10,2) AS AVG_SAL

EMPLOYEE X

Search Results Export Cost: 10ms Total 0

	department	avg_sal
--	------------	---------

22 ▶ Run | +Tab | JSON | ☰Select | ☱Ask AI

✓ 23 SELECT EMP_ID, EMP_NAME, SALARY

24 FROM EMPLOYEE

25 GROUP BY EMP_ID

26 HAVING SALARY>20000; 6ms

27 ▶ Run | +Tab | JSON | ☰Select | ☱Ask AI

28 SELECT DEPARTMENT, AVG(SALARY)::NUMERIC(10,2) AS AVG_SAL

29 FROM EMPLOYEE

EMPLOYEE X

Search Results Export Cost: 12ms Total 5

	emp_id	emp_name	salary
>	1	Aman	30000.00
>	6	Aditi	28000.00
>	7	Aanya	26000.00
>	2	Sam	25000.00
>	5	Rohan	24500.00

```
19   ▶ Run | +Tab | JSON | Ask AI
20   SELECT DEPARTMENT, AVG(SALARY)::NUMERIC(10,2) AS AVG_SAL
21   FROM EMPLOYEE
22   GROUP BY DEPARTMENT; 11ms
23
24   ▶ Run | +Tab | JSON | Select | Ask AI
25   SELECT EMP_ID, EMP_NAME, SALARY
26   FROM EMPLOYEE
27   GROUP BY EMP_ID
28   HAVING SALARY>20000;
29
30   ▶ Run | +Tab | JSON | Select | Ask AI
31   SELECT DEPARTMENT, AVG(SALARY)::NUMERIC(10,2) AS AVG_SAL
32   FROM EMPLOYEE
EMPLOYEE X
Search Results Export Cost: 75ms < 1 > Total 3


|   | department | avg_sal  |
|---|------------|----------|
| 1 | Finance    | 22250.00 |
| 2 | IT         | 27000.00 |
| 3 | HR         | 23000.00 |

  

▶ Run | Select | Ask AI
9 ✓ INSERT INTO EMPLOYEE
10 VALUES
11 (1, 'Aman', 'IT', 30000, '2023-05-23'),
12 (2, 'Sam', 'IT', 25000, '2016-05-23'),
13 (3, 'Neha', 'HR', 18000, '2025-09-19'),
14 (4, 'Suman', 'Finance', 20000, '2021-11-06'),
15 (5, 'Rohan', 'Finance', 24500, '2023-10-23'),
16 (6, 'Aditi', 'HR', 28000, '2018-04-16'),
17 (7, 'Aanya', 'IT', 26000, '2022-07-07'); 20ms AffectedRows: 7
18
▶ Run | +Tab | JSON | Ask AI
19 ✓ SELECT DEPARTMENT, AVG(SALARY)::NUMERIC(10,2) AS AVG_SAL
Result X
Search Results Export Cost: 233ms < >


Execution completed in 233ms


```

The screenshot shows a SQL editor interface with the following code:

```
Experiment-2 > Experiment-2.sql > ...
You, 13 seconds ago | 1 author (You) | Run | Select | Ask AI | Active: 127.0.0.1 | dbms_class
✓ 1 CREATE TABLE EMPLOYEE(
2     EMP_ID NUMERIC PRIMARY KEY,
3     EMP_NAME VARCHAR(20),
4     DEPARTMENT VARCHAR(20),
5     SALARY NUMERIC(10,2),
6     JOINING_DATE DATE
7 ); 233ms
8
9 INSERT INTO EMPLOYEE
10 VALUES
11 (1, 'Aman', 'IT', 30000, '2023-05-23'),
12 (2, 'Sam', 'IT', 25000, '2016-05-23'),
13 (3, 'Neha', 'HR', 18000, '2025-09-19'),
14 (4, 'Simon', 'Finance', 20000, '2021-11-05')
```

The interface includes a toolbar with icons for Run, Select, Ask AI, Export, and other database operations. Below the toolbar, a message says "Execution completed in 233ms".

6. Learning Outcome

Through this session, I have developed the following competencies:

- **Analytical Skills:** Gained the ability to transform raw row-level data into high-level summary reports using aggregation.
- **Query Logic:** Understood the logical execution order of SQL clauses: FROM → WHERE → GROUP BY → HAVING → SELECT → ORDER BY.
- **Practical Exposure:** Experienced handling real-world data scenarios, such as department-wise salary analysis and performance-based filtering in a professional database environment.